



# PLAYING MASTERMIND WITH REINFORCEMENT LEARNING



Nithish Bolleddula

Shreejaya Bharathan



# Outline


---

- Problem introduction
- Reinforcement Learning framing
- Modelling approach
- Interactive Demo
- Results
- What didn't work
- Future improvements
- Conclusion

# Introduction

---

- Mastermind is a two player code breaking game
- The aim of the game is to guess the (4 digit) code in  $n$  - turns



# Mastermind as an RL problem

---


- Environment : Mastermind game board
- Agent : Plays a move guessing the pattern
- State : Any combination of 4 colours as a guess
- Feedback : (Number of colors guessed correctly in right position, Number of colors guessed but in wrong position)
- Reward : +1, -1

# Modeling Approach

---

- Existing solution - 5 guess algorithm
- Baseline - random agent

We implemented two RL frameworks:

- Q learning - Final model 
- Policy gradient - Another RL approach

# Demo of our final model!

---

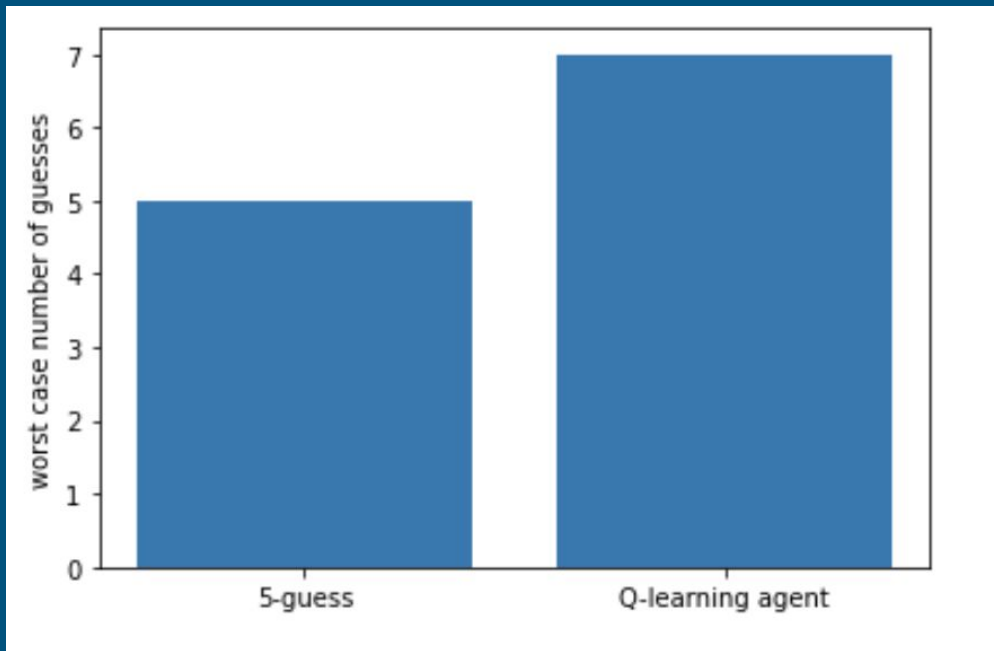
Interactive demo - [https://github.com/ShreejayaB/mastermind/blob/master/Interactive\\_play.ipynb](https://github.com/ShreejayaB/mastermind/blob/master/Interactive_play.ipynb)

Q learning code - <https://github.com/ShreejayaB/mastermind/blob/master/Q-learning.ipynb>

# Results

---

Q learning agent is able to perform close to the popular 5 guess algorithm by Donald Knuth.



# How to interpret our results?

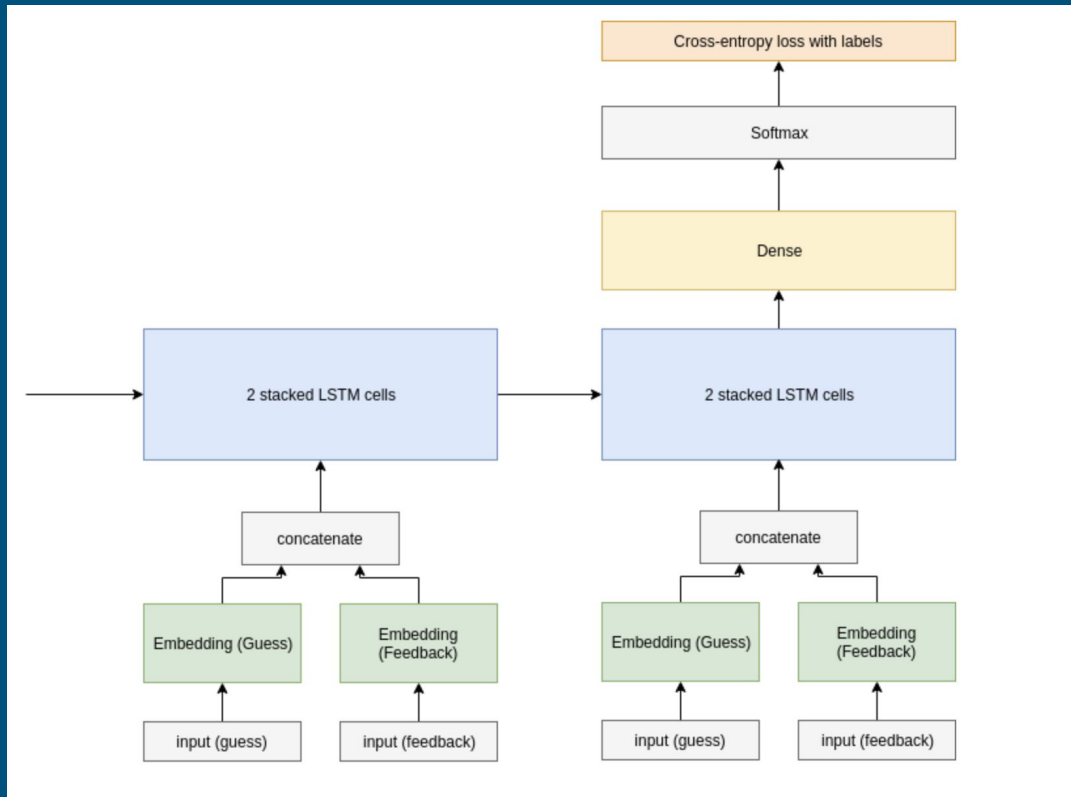
---

- Initially Q learning was unstable
- After restricting the state space, we obtained a stable algorithm
- The reason for instability is the complex nature of the game - environment, feedback and reward change with each episode based on secret code
- Given the changing nature of the environment, implementing an RL algorithm to outperform the deterministic 5 guess is challenging



# Policy Gradient approach & what didn't work

- To learn different game patterns, we used an LSTM network
- The LSTM learns embeddings of the guesses
- This meant more data simulation, making the process VERY slow



# Future improvements

---

- Parallelize the learning process
- Use GRU instead of LSTM
- Simulate the data efficiently
- Decrease the size of embedding of state and action vectors

# Conclusions

---

- Mastermind is a fairly complex reinforcement learning problem
- Q Learning becomes unstable for dynamic environments
- Double Q learning may also help improve performance
- Policy gradients are computationally intensive. Use GPUs!!
- It's not so easy to win mastermind afterall!



Thank you!

